

CSC 411

Computer Organization (Fall 2024)
Lecture 1: Logistics

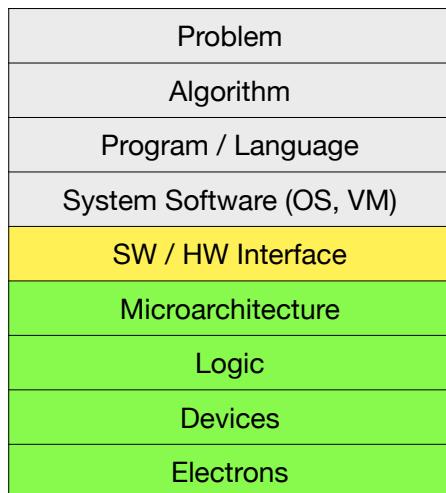
Prof. Marco Alvarez, University of Rhode Island

A computing system

ISA (Instruction Set Architecture)

Implementation of the ISA

Building blocks (gates)



What is this course about?

- A comprehensive exploration of fundamental principles in computer organization
 - the fascinating interplay between hardware and software that powers modern computing systems, from transistors to high-level programming
- Topics covered
 - instruction sets, assembly language, processor design, the memory hierarchy, and performance optimization, with a particular focus on the **RISC-V architecture**

How computers work?

Computer architecture

- What is computer architecture
 - the science of designing and implementing computing systems
 - primary emphasis on the HW/SW interface, particularly the instruction set architecture, processor design, and memory systems
- Key design goals
 - maximizing performance (execute tasks as quickly as possible)
 - optimizing energy efficiency (while still meeting performance)
 - balancing cost and performance
- Design priorities vary by system type
 - e.g., supercomputers emphasize throughput, smartphones prioritize energy efficiency
 - despite varying priorities, **fundamental principles** apply across different system architectures

Why study computer architecture?

- Understand current/future capabilities of computing systems
 - why computers work as they do (principles and design decisions)
- Develop better software
 - best system programmers understand all abstraction levels and the underlying hardware
 - write more efficient code, debug complex issues, exploit hardware characteristics
- Optimizing performance across the entire stack
 - writing well tuned software by understanding the relationship between hardware and software
- Establish foundation for further work
 - prepare for advanced study and career opportunities

To illustrate the potential gains from performance engineering, consider multiplying two 4096-by-4096 matrices. Here is the four-line kernel of Python code for matrix-multiplication:

```
for i in xrange(4096):
    for j in xrange(4096):
        for k in xrange(4096):
            C[i][j] += A[i][k] * B[k][j]
```

Version	Implementation	Running time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak (%)
1	Python	25,552.48	0.005	1	—	0.00
2	Java	2,372.68	0.058	11	10.8	0.01
3	C	542.67	0.253	47	4.4	0.03
4	Parallel loops	69.80	1.969	366	7.8	0.24
5	Parallel divide and conquer	3.80	36.180	6,727	18.4	4.33
6	plus vectorization	1.10	124.914	23,224	3.5	14.96
7	plus AVX intrinsics	0.41	337.812	62,806	2.7	40.45

From: "There's plenty of room at the Top: What will drive computer performance after Moore's law?"

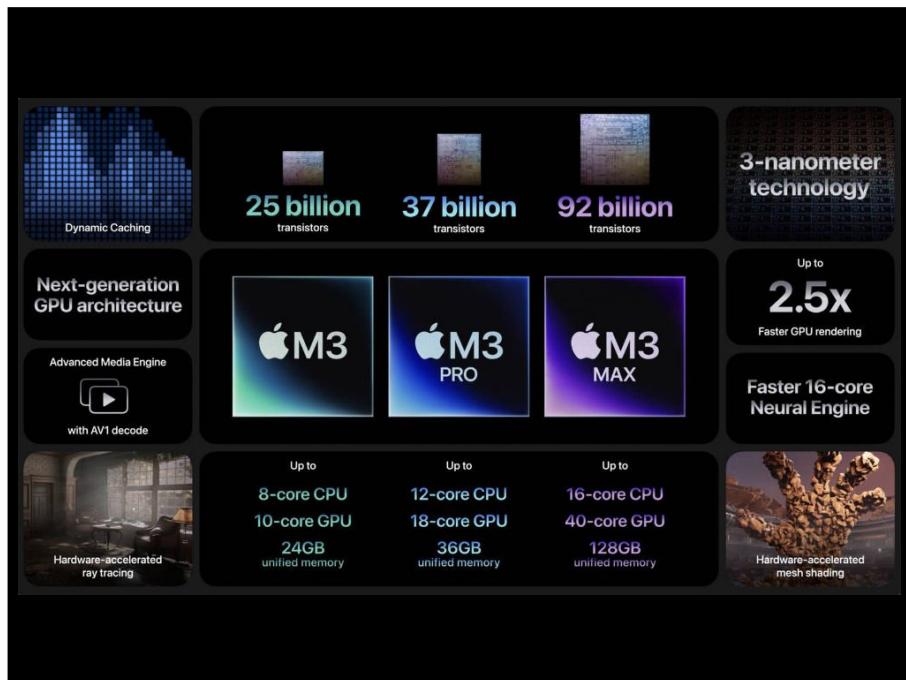
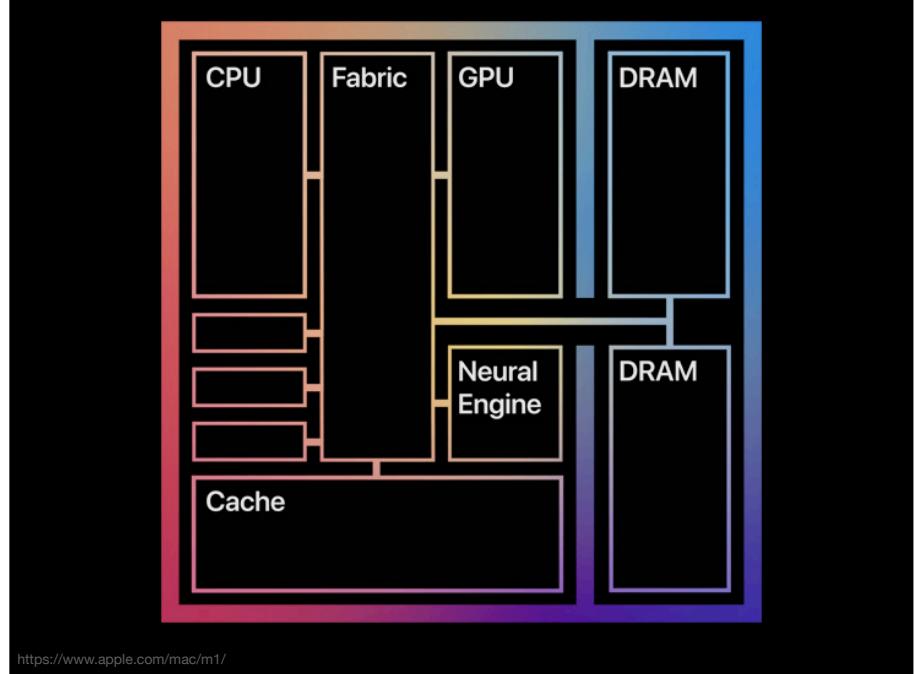
Why learning Assembly?

- Critical for understanding machine-level execution
 - provides a low-level perspective on how computers execute instructions
- Key advantages
 - enhanced debugging skills
 - optimizing program performance and identifying sources of program inefficiency
 - implementing system software that leverages hardware capabilities
 - enhanced security against malware

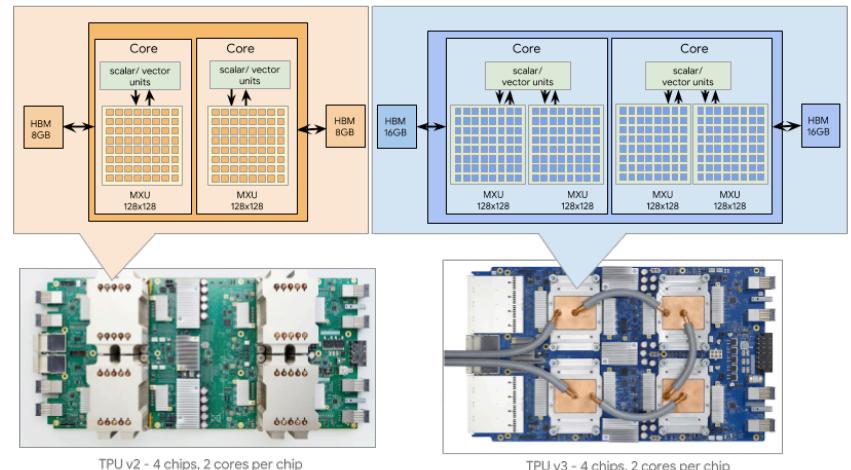
Modern computer architecture

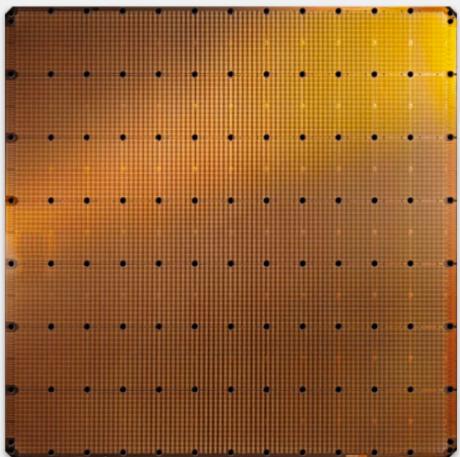
- Key strategies for higher performance and efficiency
 - **co-design** across the hierarchy (bring algorithms to devices)
 - **specialize** as much as possible for targeted efficiency
- Outlook
 - fundamental building blocks and principles persist
 - emergence of novel architectures
 - evolving and powerful landscape
 - exciting era for computer architecture

Examples



TPU (Tensor Processing Unit)





Cerebras WSE-2
46,225mm² Silicon
2.6 Trillion transistors

Largest GPU
826mm² Silicon
54.2 Billion transistors

<https://www.cerebras.net/product-chip/>

Cerebras's Wafer Scale ML Engine

Course organization

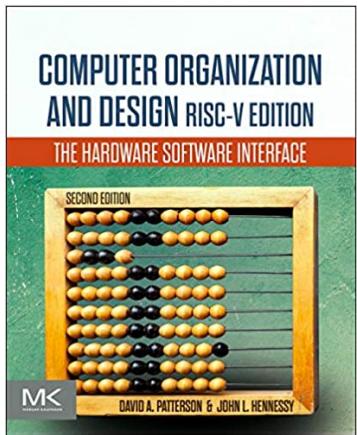
Course website

- [URL](#)
 - <https://homepage.cs.uri.edu/~malvarez/teaching/csc-411>
 - Syllabus
 - Schedule
 - Presentations
 - Resources

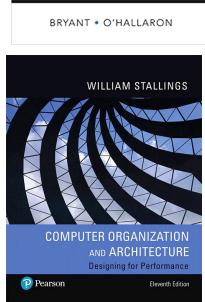
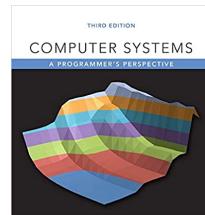
Course information

- [Lectures](#)
 - MW 4:30-5:45p
- [Lab](#)
 - T 5-5:50p (via Zoom)
- [Team](#)
 - Marco Alvarez, Instructor
 - Emily Light
 - Liam Cannon
- [Office Hours](#)
 - TBA

Recommended textbooks



Required



Digital Design and
Computer Architecture
RISC-V Edition



Support tools



Ed Discussion: Academic discussion, polls, quizzes.



Gradescope: Assignment submission and grading.



Zoom: Virtual labs and office hours.

Grading

- Homework assignments (15%)
- Midterm exam (30%)
- Technical presentation (20%)
- Final exam (35%)

Coursework

- Homework assignments
 - problem sets and programming tasks (primarily C and RISC-V)
 - individual work, however discussions and collaboration are allowed
 - you **must write your own** code and solutions
 - late submissions **NOT accepted**
 - ample time given to complete (6-9 days)
 - start early and use office hours for guidance and feedback
- Exams
 - in-person and open-book (printed materials only)
 - no electronic devices allowed
 - mix of multiple-choice, short answer, and problem-solving questions
 - designed to test understanding, not just memorization

Coursework

› Technical presentation

- team work (3 members), diving deep into an advanced topic
 - specific topics provided by instructor, covering cutting-edge areas in computer organization
 - prepare 20-minute presentation to share your insights/findings
- enhance your communication and collaboration skills
- contribute to the class' collective knowledge
- **outstanding presentations** will receive extra credit

Academic integrity

› Assignments

- collaboration is encouraged for discussing ideas, but each student/team must submit their own unique solutions
- sharing/copying solutions from peers is **prohibited**
- using uncredited AI-generated content **prohibited**

› AI and LLMs

- AI tools (e.g., ChatGPT, Gemini, Claude, GitHub Copilot) can be used to enhance learning through brainstorming, concept exploration, and strategy development
 - students must critically evaluate and fully understand any AI-generated content used in their work
 - all AI-assisted work **must be cited** in submissions
- AI tools are designed to support students' learning, NOT to replace independent problem-solving and critical thinking
- seek instructor guidance if uncertain about appropriate AI tool usage

How to succeed?

› Class attendance

- students are expected to **attend all lectures and labs**, which are held **synchronously** and will not be recorded
- regular attendance is linked to higher grades and better comprehension of course material

› Active participation

- students should engage critically by asking questions, participating in discussions, and thinking deeply about the material
 - during online labs, cameras should be turned on
- **laptops and cellphones are NOT permitted** unless being used for note-taking
- students are advised to utilize Ed and attend office hours for additional support

Aim for excellence in your work, rather than simply chasing a good grade