

# CSC 411

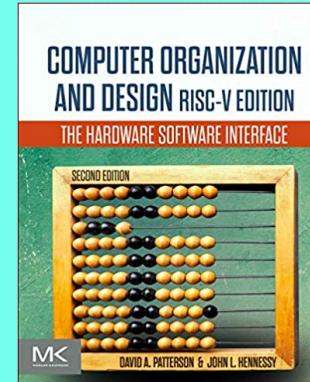
Computer Organization (Fall 2024)  
Lecture 9: Computer systems

Prof. Marco Alvarez, University of Rhode Island

## Units of measure

## Disclaimer

Some figures and slides are adapted from:  
Computer Organization and Design (Patterson and Hennessy)  
The Hardware/Software Interface



## SI prefixes

Factor	Name	Symbol	Factor	Name	Symbol
$10^1$	deca	da	$10^{-1}$	deci	d
$10^2$	hecto	h	$10^{-2}$	centi	c
$10^3$	kilo	k	$10^{-3}$	milli	m
$10^6$	mega	M	$10^{-6}$	micro	$\mu$
$10^9$	giga	G	$10^{-9}$	nano	n
$10^{12}$	tera	T	$10^{-12}$	pico	p
$10^{15}$	peta	P	$10^{-15}$	femto	f
$10^{18}$	exa	E	$10^{-18}$	atto	a
$10^{21}$	zetta	Z	$10^{-21}$	zepto	z
$10^{24}$	yotta	Y	$10^{-24}$	yocto	y
$10^{27}$	ronna	R	$10^{-27}$	ronto	r
$10^{30}$	quette	Q	$10^{-30}$	quecto	q

SI prefixes are a set of 24 prefixes used in the International System of Units (SI) to indicate multiples and submultiples of SI units. They are based on powers of 10, and each prefix has a unique symbol.

## Binary prefixes

The SI prefixes refer strictly to powers of 10. They should not be used to indicate powers of 2 (for example, one kilobit represents 1000 bits and not 1024 bits). The names and symbols for prefixes to be used with powers of 2 are recommended as follows:

kibi	Ki	$2^{10}$
mebi	Mi	$2^{20}$
gibi	Gi	$2^{30}$
tebi	Ti	$2^{40}$
pebi	Pi	$2^{50}$
exbi	Ei	$2^{60}$
zebi	Zi	$2^{70}$
yobi	Yi	$2^{80}$

Decimal term	Abbreviation	Value	Binary term	Abbreviation	Value	% Larger
kilobyte	KB	$10^3$	kibibyte	KiB	$2^{10}$	2%
megabyte	MB	$10^6$	mebibyte	MiB	$2^{20}$	5%
gigabyte	GB	$10^9$	gibibyte	GiB	$2^{30}$	7%
terabyte	TB	$10^{12}$	tebibyte	TiB	$2^{40}$	10%
petabyte	PB	$10^{15}$	pebibyte	PiB	$2^{50}$	13%
exabyte	EB	$10^{18}$	exbibyte	EiB	$2^{60}$	15%
zettabyte	ZB	$10^{21}$	zebibyte	ZiB	$2^{70}$	18%
yottabyte	YB	$10^{24}$	yobibyte	YiB	$2^{80}$	21%
ronnabyte	RB	$10^{27}$	robibyte	RiB	$2^{90}$	24%
queccabyte	QB	$10^{30}$	quebibyte	QiB	$2^{100}$	27%

# Computer systems

## The computer revolution

- Progress in computer technology
  - underpinned by domain-specific accelerators
- Novel applications (not long ago considered fiction):
  - generative AI, automotive computers, smartphones, human genome project, web, search engines
- Although common hardware technologies ...
  - different design requirements

Computers in all forms are now pervasive

## Classes of computers

- Personal computers (desktops/laptops)
  - low cost, general purpose, variety of software
  - subject to cost/performance tradeoff
- Server computers
  - network based, large workloads
  - high capacity, performance, reliability
  - range from small servers to building sized supercomputers (high-end scientific and engineering calculations)
- Embedded computers
  - largest class by volume
  - hidden as components of systems (Internet of Things)
  - power/performance/cost constraints

## The post PC era

- Personal Mobile Device (PMD) — overtaking PCs
  - battery operated, internet-connected
  - smart phones, tablets, electronic glasses
- Cloud computing — overtaking traditional servers
  - Warehouse Scale Computers (WSC) — e.g. Amazon AWS
  - Software as a Service (SaaS)
    - portion of software run on a PMD and a portion run in the Cloud

## Accelerators and heterogeneous computing

- Multiple cores in one chip
- System-on-a-Chip (SoC) design
- Graphics processing units (GPUs)
  - throughput-oriented multicore processors
  - great for gaming and machine learning
- Focus on performance and energy efficiency
  - critical in all classes of computers

High performance computing

## Top 500 project

- 500 most powerful computers in the world
  - exascale applications
- Updated twice a year
  - ISC'xy in June, Germany and SC'xy in November, U.S.

NOVEMBER 2023						
	Manufacturer	Computer	Country	Cores	Rmax [PFLOPS]	Power [MW]
1	Oak Ridge National Laboratory	HPE	Frontier	8,730,112	1,102	21.1
2	Argonne National Laboratory	HPE	HPE Cray EX235a, AMD EPYC 64C 2.6GHz, Intel Mi250X, Slingshot-11	4,742,808	585.3	24.6
3	Microsoft Azure	Microsoft	Eagle	1,123,200	561.2	
4	RIKEN Center for Computational Science	Fujitsu	Fugaku	7,630,848	442.0	29.9
5	EuroHPC / CSC	HPE	HPE Cray EX235a, AMD EPYC 64C 2.6GHz, Intel Mi250X, Slingshot-11	2,069,760	309.1	6.0

credit: Berkeley's CS267 lectures

## Exascale applications at LBNL



credit: Berkeley's CS267 lectures

## Top computer in the world



### Frontier (#1) System Overview



#### System Performance

- Peak performance of 1.6 double precision exaFLOPS
- Measured Top500 performance (Rmax) was 1.102 exaFLOPS

#### Each node has

- 3rd Gen AMD EPYC CPU with 64 cores
- 4 Purpose Built AMD Instinct 250X GPUs
- 4X128 GB of fast memory, 1 per GPU
- 5 terabytes of flash memory



credit: Berkeley's CS267 lectures

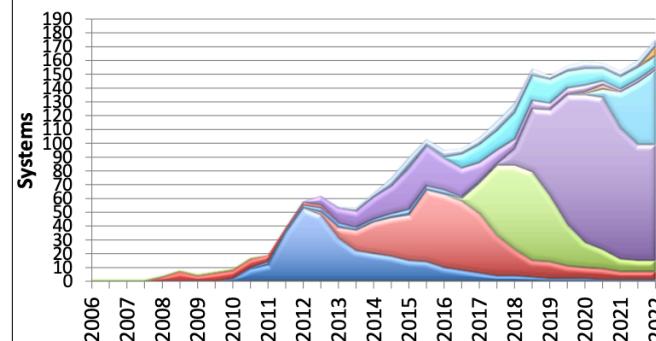
#### The system includes

- 9,472 nodes
- Slingshot interconnect

## Heterogenous devices



- Nvidia Fermi
- Nvidia Kepler
- Nvidia Pascal
- Nvidia Volta
- Nvidia Ampere
- Nvidia Turing
- ATI Radeon
- IBM Cell
- Clearspeed
- Intel Xeon Phi
- Xeon Phi Main
- AMD Instinct
- Others



2022: Over 1/3<sup>rd</sup> of all systems have accelerators or co-processors

The actual "performance share" is higher

credit: Berkeley's CS267 lectures

# Abstractions

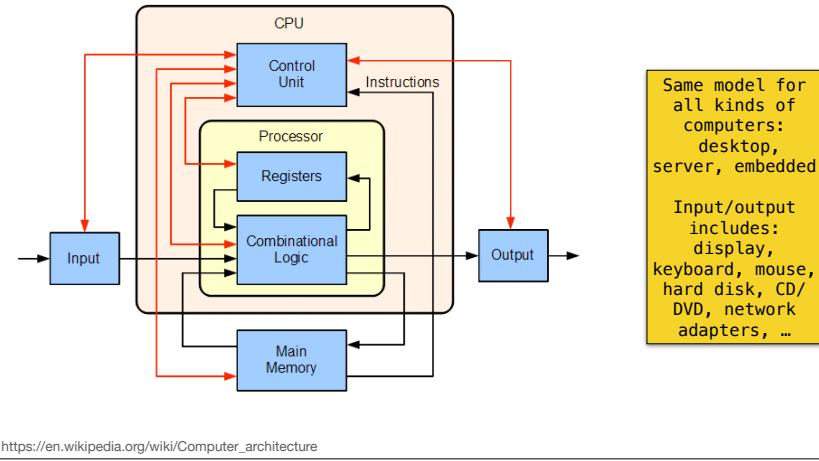
## Seven great ideas

- Use **abstraction** to simplify design
- Make the **common case fast**
- Performance via **parallelism**
- Performance via **pipelining**
- Performance via **prediction**
- **Hierarchy** of memories
- **Dependability** via redundancy

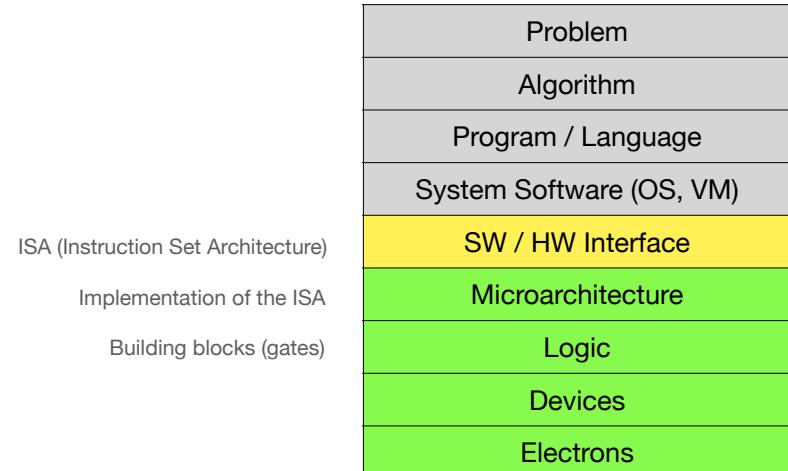


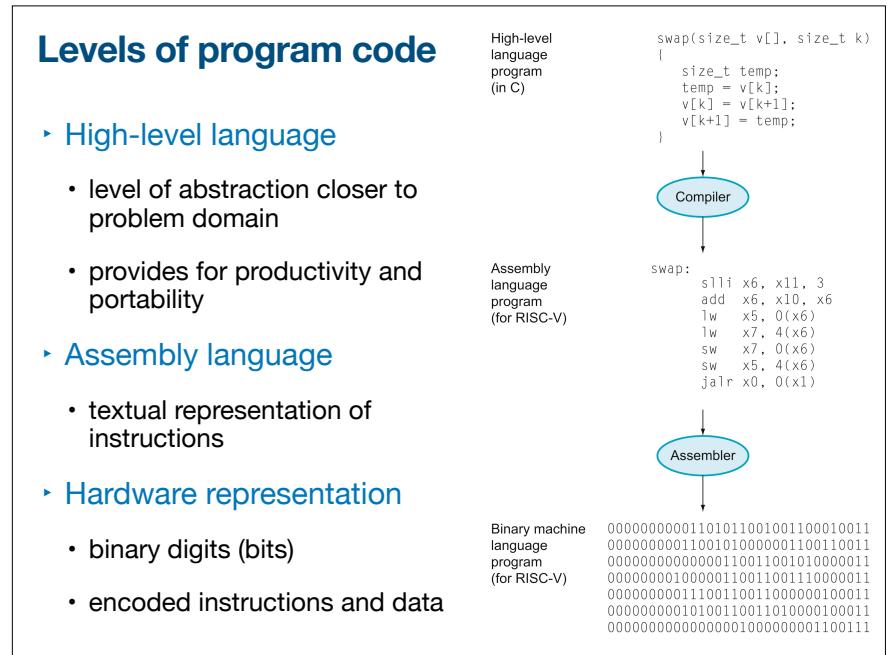
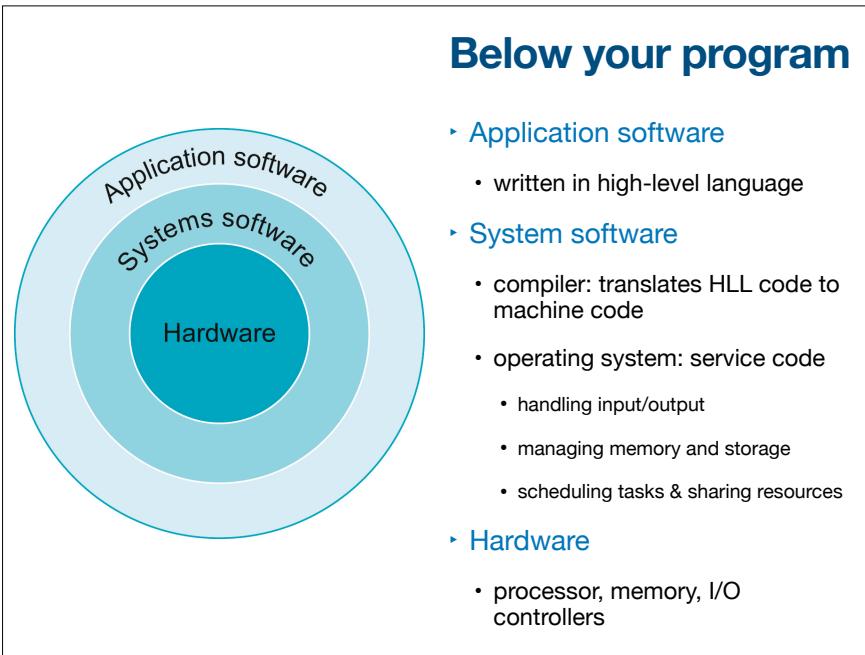
## The Von-Neumann model

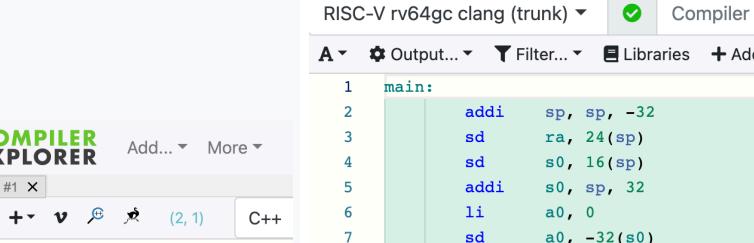
- High-level organization of computer hardware



## Abstraction layers (computing system)







The screenshot shows the Compiler Explorer interface with the following details:

- Compiler:** RISC-V rv64gc clang (trunk)
- Source Code (C++ source #1):**

```
1 #include <stdio.h>
2
3 int main() {
4     printf("hello world !\n");
5
6     return 0;
7 }
```
- Assembly Output:**

```
1 main:
2         addi    sp, sp, -32
3         sd     ra, 24(sp)
4         sd     s0, 16(sp)
5         addi   s0, sp, 32
6         li     a0, 0
7         sd     a0, -32(s0)
8         sw     a0, -20(s0)
9         lui    a0, %hi(.L.str)
10        addi   a0, a0, %lo(.L.str)
11        call   printf
12        ld     a0, -32(s0)
13        ld     ra, 24(sp)
14        ld     s0, 16(sp)
15        addi   sp, sp, 32
16        ret
17 .L.str:
18         .asciz "hello world !\n"
```

# Looking into an executable

```
#include <stdio.h>

int main() {
    printf("hello world !\n");

    return 0;
}

$ objdump -d
```

# Instruction set architecture

"the words of a computer's language are called *instructions*, and its vocabulary is called an *instruction set*" [Computer Organization and Design, P&H]

the basic job of a CPU is to execute instructions

## Major ISAs

### Major ISAs

- X86 (Intel and AMD)
  - desktop computers, laptops, servers
- ARM
  - widely used on embedded, phones, recently on laptops, servers
- RISC-V
  - fast-growing, used in embedded systems, servers, personal computers

### Even more microarchitectures

- Apple/Samsung/Qualcomm have their own microarchitecture (implementation) of ARM
- Intel/AMD have different microarchitectures for x86

## Definitions

### Instruction set

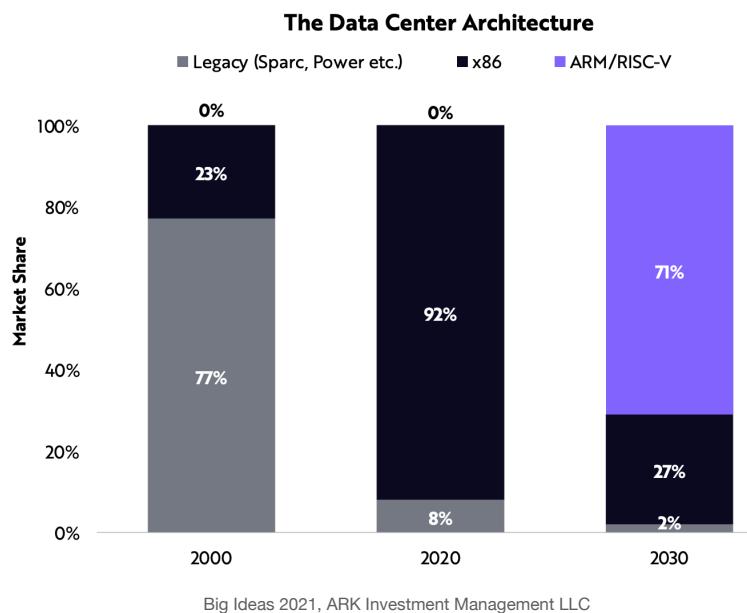
- the **repertoire** of instructions of a computer
- different computers have different instruction sets
  - ... with many aspects in common

### Instruction set architecture (ISA)

- a **contract** between hardware and software
- defines instruction format, syntax, and semantics

### Microarchitecture

- implementation of the architecture, i.e., how the processor physically executes the instructions
- e.g., cache sizes, core frequency, pipelining, branch prediction
- different microarchitectures for the same ISA (e.g., Intel vs AMD for x86)



## CISC vs RISC

### ‣ Complex Instruction Set Computer (CISC)

- more complex instructions, can directly manipulate memory
- X86, X86\_64 (Intel and AMD, desktop/laptop/server)
  - X86\* internally are still RISC

### ‣ Reduced Instruction Set Computer (RISC)

- simpler instructions, load-store architecture
- ARM (smartphone/pad)
- **RISC-V** (free ISA, closer to MIPS than other ISAs)
- Others: Power ISA, SPARC, etc



ACM has named **John L. Hennessy**, former President of Stanford University, and **David A. Patterson**, retired Professor of the University of California, Berkeley, recipients of the **2017 ACM A.M. Turing Award** for pioneering a systematic, quantitative approach to the design and evaluation of computer architectures with enduring impact on the microprocessor industry.

<https://www.youtube.com/watch?v=3LVeEjsn8Ts>

## CISC vs RISC

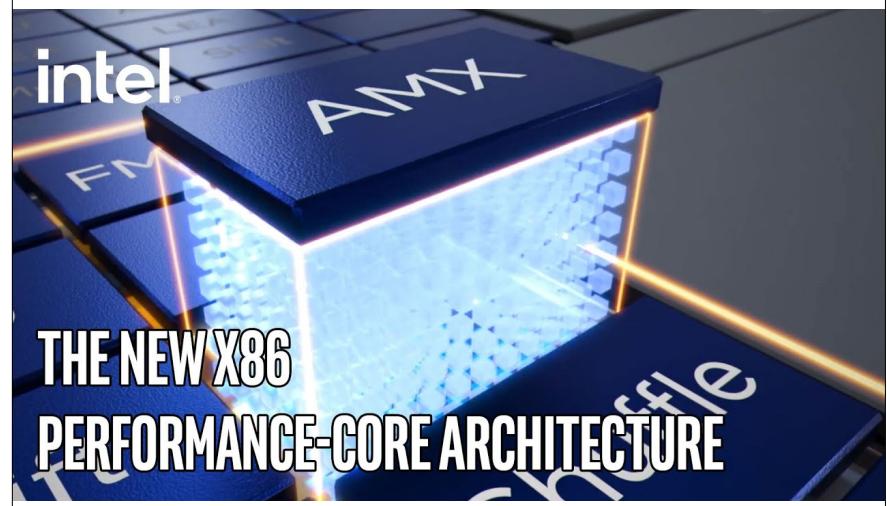
### ‣ CISC

- early trend was to add more and more instructions
- VAX architecture had an instruction to multiply polynomials!
- operations may directly manipulate memory

### ‣ RISC philosophy

- keep the instruction set **small and simple**
- makes it easier to build fast hardware.
- let software do complicated operations by composing simpler ones

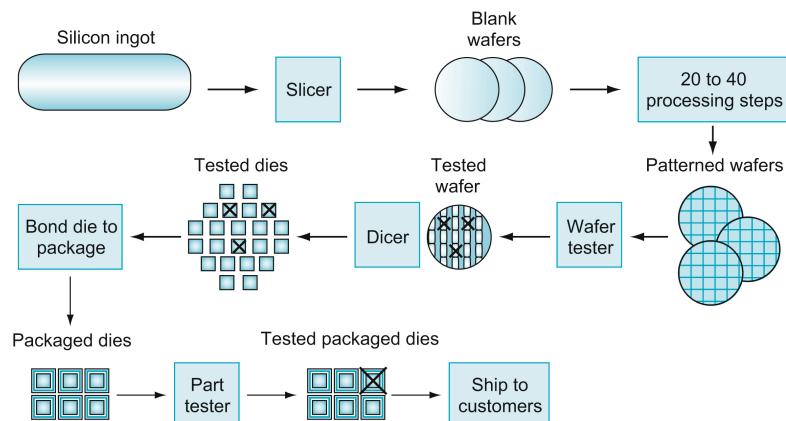
## Intel's x86 core



<https://www.youtube.com/watch?v=ijTRvIQV7bE>

# Chip Manufacturing

## Manufacturing integrated circuits



Yield: proportion of working dies per wafer

## Semiconductor technology

### Conductors

- materials that easily conduct



### Insulators

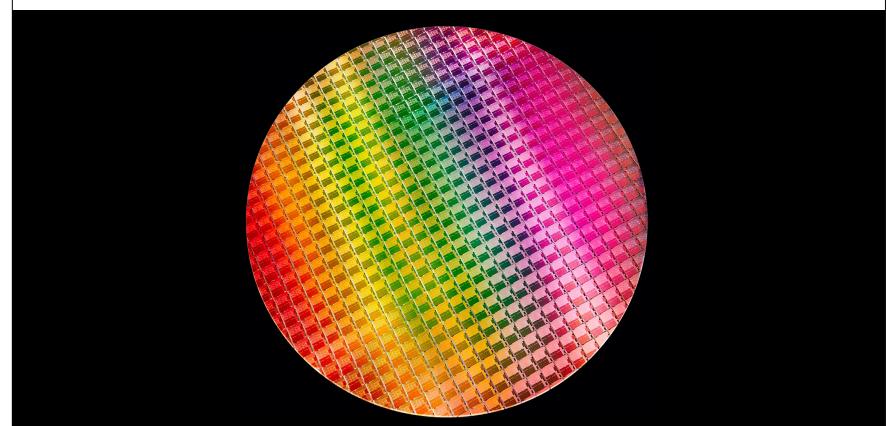
- materials that do not conduct

### Semiconductors

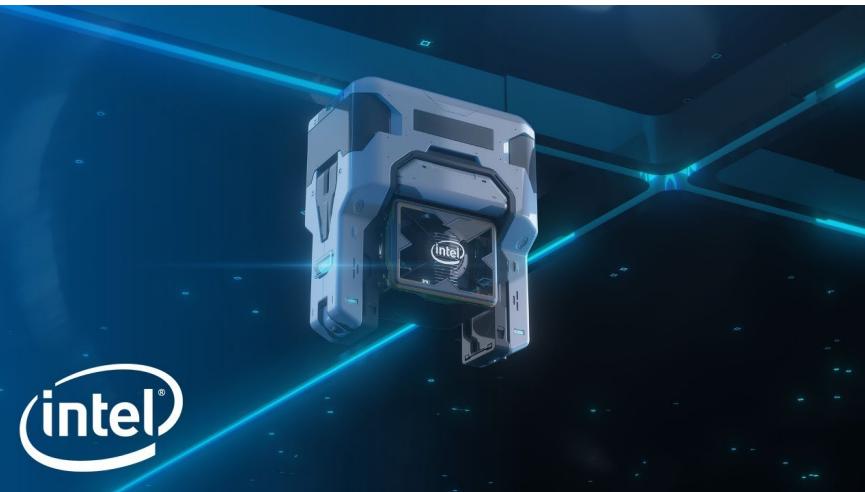
- materials with conductivity between conductors and insulators
- those properties can be controlled by doping (adding/removing electrons)
- Silicon (base material for computer technology)

## Intel® Core 10th Gen

- 300mm wafer (almost 1 foot), 506 chips, 10nm technology
- Each chip is 11.4 x 10.7 mm (a bit less than .5 in)



## From sand to silicon



[https://www.youtube.com/watch?v= VMYPLXnd7E](https://www.youtube.com/watch?v=VMYPLXnd7E)