

# CSC 411

Computer Organization (Fall 2024)  
Lecture 9: Computer systems

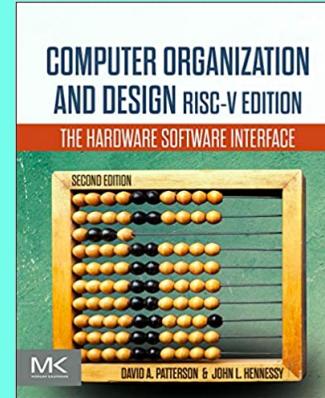
Prof. Marco Alvarez, University of Rhode Island

## Units of measure

## Disclaimer

Some figures and slides are adapted from:

Computer Organization and Design (Patterson and Hennessy)  
The Hardware/Software Interface



## SI prefixes

Factor	Name	Symbol	Factor	Name	Symbol
$10^1$	deca	da	$10^{-1}$	deci	d
$10^2$	hecto	h	$10^{-2}$	centi	c
$10^3$	kilo	k	$10^{-3}$	milli	m
$10^6$	mega	M	$10^{-6}$	micro	$\mu$
$10^9$	giga	G	$10^{-9}$	nano	n
$10^{12}$	tera	T	$10^{-12}$	pico	p
$10^{15}$	peta	P	$10^{-15}$	femto	f
$10^{18}$	exa	E	$10^{-18}$	atto	a
$10^{21}$	zetta	Z	$10^{-21}$	zepto	z
$10^{24}$	yotta	Y	$10^{-24}$	yocto	y
$10^{27}$	ronna	R	$10^{-27}$	ronto	r
$10^{30}$	quette	Q	$10^{-30}$	quecto	q

SI prefixes are a set of 24 prefixes used in the International System of Units (SI) to indicate multiples and submultiples of SI units. They are based on powers of 10, and each prefix has a unique symbol.

## Binary prefixes

The SI prefixes refer strictly to powers of 10. They should not be used to indicate powers of 2 (for example, one kilobit represents 1000 bits and not 1024 bits). The names and symbols for prefixes to be used with powers of 2 are recommended as follows:

kibi	Ki	$2^{10}$
mebi	Mi	$2^{20}$
gibi	Gi	$2^{30}$
tebi	Ti	$2^{40}$
pebi	Pi	$2^{50}$
exbi	Ei	$2^{60}$
zebi	Zi	$2^{70}$
yobi	Yi	$2^{80}$

Decimal term	Abbreviation	Value	Binary term	Abbreviation	Value	% Larger
kilobyte	KB	$10^3$	kibibyte	KiB	$2^{10}$	2%
megabyte	MB	$10^6$	mebibyte	MiB	$2^{20}$	5%
gigabyte	GB	$10^9$	gibibyte	GiB	$2^{30}$	7%
terabyte	TB	$10^{12}$	tebibyte	TiB	$2^{40}$	10%
petabyte	PB	$10^{15}$	pebibyte	PiB	$2^{50}$	13%
exabyte	EB	$10^{18}$	exbibyte	EiB	$2^{60}$	15%
zettabyte	ZB	$10^{21}$	zebibyte	ZiB	$2^{70}$	18%
yottabyte	YB	$10^{24}$	yobibyte	YiB	$2^{80}$	21%
ronnabyte	RB	$10^{27}$	robibyte	RiB	$2^{90}$	24%
queccabyte	QB	$10^{30}$	quebabyte	QiB	$2^{100}$	27%

## Practice

- A hard drive is advertised as 1 TB (terabyte)
  - calculate the difference in storage capacity (in bytes) between the SI and binary interpretations of this size
- A file size is given as 2.5 GiB (gibibytes)
  - express this in both gigabytes (GB) and mebibytes (MiB)

## Computer systems

## The computer revolution

- Progress in computer technology
  - underpinned by domain-specific accelerators
- Novel applications (not long ago considered fiction):
  - generative AI, automotive computers, smartphones, human genome project, web, search engines
- Although common hardware technologies ...
  - different design requirements

Computers in all forms are now pervasive

## Classes of computers

- Personal computers (desktops/laptops)
  - low cost, general purpose, variety of software
  - subject to cost/performance tradeoff
- Server computers
  - network based, large workloads
  - high capacity, performance, reliability
  - range from small servers to building sized supercomputers (high-end scientific and engineering calculations)
- Embedded computers
  - largest class by volume
  - hidden as components of systems (Internet of Things)
  - power/performance/cost constraints

## The post PC era

- Personal Mobile Device (PMD) — overtaking PCs
  - battery operated, internet-connected
  - smart phones, tablets, electronic glasses
- Cloud computing — overtaking traditional servers
  - Warehouse Scale Computers (WSC) — e.g. Amazon AWS
  - Software as a Service (SaaS)
    - portion of software run on a PMD and a portion run in the Cloud

## Accelerators and heterogeneous computing

- Multiple cores in one chip
- System-on-a-Chip (SoC) design
- Graphics processing units (GPUs)
  - throughput-oriented multicore processors
  - great for gaming and machine learning

Focus on performance and energy efficiency  
critical in all classes of computers

# High performance computing

## Top 500 project

- 500 most powerful computers in the world
  - exascale applications
- Updated twice a year
  - ISC'xy in June, Germany and SC'xy in November, U.S.

Rank	Organization	Manufacturer	Computer	Country	Cores	Rmax [PFLOPS]	Power [MW]
1	Oak Ridge National Laboratory	HPE	Frontier HPE Cray EX235a, AMD EPYC 64C 2.0GHz, Instinct MI250X, Slingshot-11	USA	8,730,112	1,102	21.1
2	Argonne National Laboratory	HPE	Aurora* HPE Cray EX Xeon CPU Max 9470 52C 2.4GHz, Intel Data Center GPU Max, Slingshot-11	USA	4,742,808	585.3	24.6
3	Microsoft Azure	Microsoft	Eagle Microsoft NDv5 Xeon Platinum 8480C 48C 2.6GHz, NVIDIA H100, NVIDIA Infiniband NDR	USA	1,123,200	561.2	
4	RIKEN Center for Computational Science	Fujitsu	Fugaku Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D	Japan	7,630,848	442.0	29.9
5	EuroHPC / CSC	HPE	LUMI HPE Cray EX235a, AMD EPYC 64C 2.0GHz, Instinct MI250X, Slingshot-11	Finland	2,069,760	309.1	6.0

credit: Berkeley's CS267 lectures

## Top computer in the world



### Frontier (#1) System Overview



#### System Performance

- Peak performance of 1.6 double precision exaFLOPS
- Measured Top500 performance (Rmax) was 1.102 exaFLOPS

#### Each node has

- 3rd Gen AMD EPYC CPU with 64 cores
- 4 Purpose Built AMD Instinct 250X GPUs
- 4X128 GB of fast memory, 1 per GPU
- 5 terabytes of flash memory

#### The system includes

- 9,472 nodes
- Slingshot interconnect



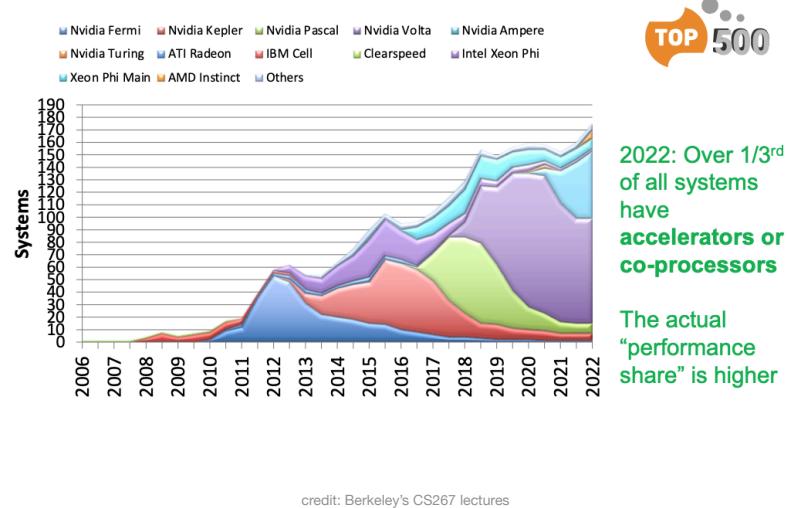
credit: Berkeley's CS267 lectures

## Exascale applications at LBNL



credit: Berkeley's CS267 lectures

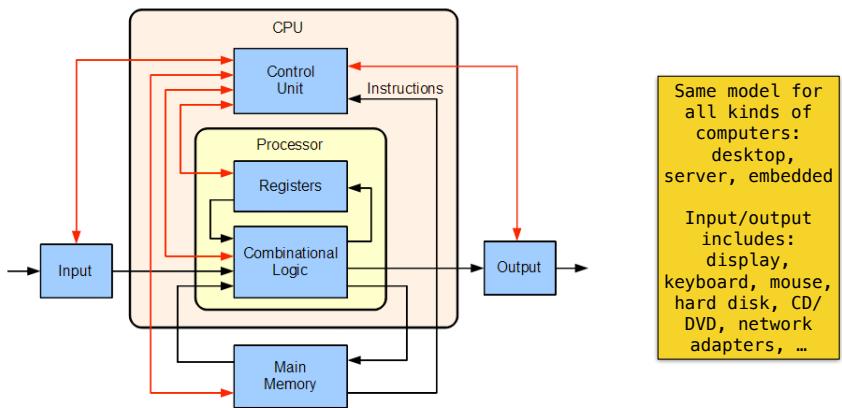
## Heterogenous devices



## Abstractions

## The Von-Neumann model

- High-level organization of computer hardware

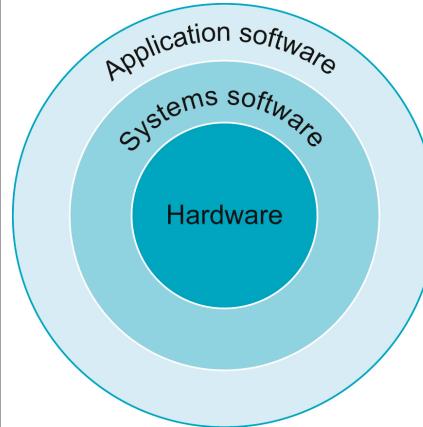
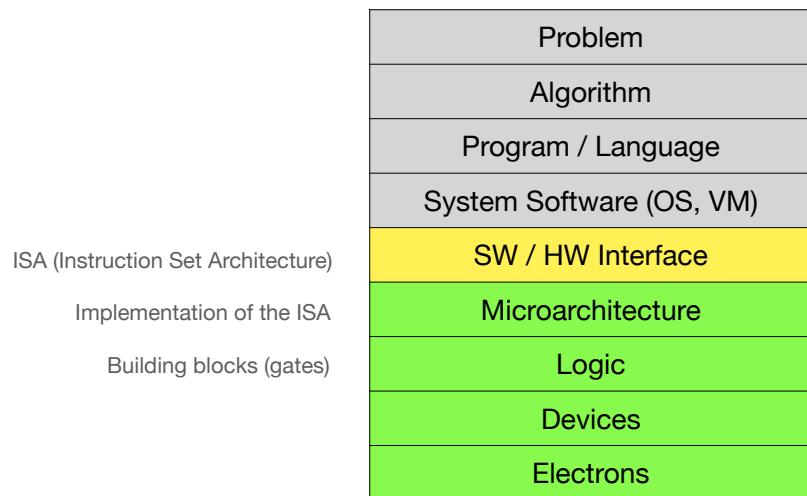


## Seven great ideas

- Use **abstraction** to simplify design
- Make the **common case fast**
- Performance via **parallelism**
- Performance via **pipelining**
- Performance via **prediction**
- Hierarchy** of memories
- Dependability** via redundancy



# Abstraction layers (computing system)



# Below your program

- Application software
  - written in high-level language
- System software
  - compiler: translates HLL code to machine code
  - operating system: service code
    - handling input/output
    - managing memory and storage
    - scheduling tasks & sharing resources
- Hardware
  - processor, memory, I/O controllers

# Levels of program code

## ‣ High-level language

- level of abstraction closer to problem domain
- provides for productivity and portability

High-level language program (in C)

```
swap(size_t v[], size_t k)
{
    size_t temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```



Assembly language program (for RISC-V)

```
slli x6, x11, 3
add x6, x10, x6
lw x5, 0(x6)
lw x7, 4(x6)
sw x7, 0(x6)
sw x5, 4(x6)
jalr x0, 0(x1)
```



## ‣ Assembly language

- textual representation of instructions

## ‣ Hardware representation

- binary digits (bits)
- encoded instructions and data

Binary machine language program (for RISC-V)

```
00000000001101011001001100010011
00000000001100101000001100110011
00000000000000110011001010000011
00000000010000011001100110000011
00000000001110011001100000100011
000000000010100110011010000100011
0000000000000000100000001100111
```

```
RISC-V rv64gc clang (trunk) ▾ ✓ Compiler op
A ▾ Output... ▾ Filter... ▾ Libraries + Add n
1 main:
2     addi   sp, sp, -32
3     sd    ra, 24(sp)
4     sd    s0, 16(sp)
5     addi  s0, sp, 32
6     li    a0, 0
7     sd    a0, -32(s0)
8     sw    a0, -20(s0)
9     lui   a0, %hi(.L.str)
10    addi  a0, a0, %lo(.L.str)
11    call  printf
12    ld    a0, -32(s0)
13    ld    ra, 24(sp)
14    ld    s0, 16(sp)
15    addi sp, sp, 32
16    ret
17 .L.str:
18    .asciz "hello world !\n"
```

<https://godbolt.org/>

## Looking into an executable

```
#include <stdio.h>

int main() {
    printf("hello world !\n");

    return 0;
}
```

```
$ objdump -d
```

## Definitions

- Instruction set
  - the repertoire of instructions of a computer
  - different computers have different instruction sets
    - ... with many aspects in common
- Instruction set architecture (ISA)
  - a contract between hardware and software
  - defines instruction format, syntax, and semantics
- Microarchitecture
  - implementation of the architecture, i.e., how the processor physically executes the instructions
  - e.g., cache sizes, core frequency, pipelining, branch prediction
  - different microarchitectures can use the same ISA (e.g., Intel vs AMD for x86)

## Instruction set architecture

“the words of a computer’s language are called *instructions*, and its vocabulary is called an *instruction set*” [Computer Organization and Design, P&H]

the basic job of a CPU is to execute instructions

## Major ISAs

- Major ISAs
  - X86 (Intel/AMD)
    - desktop computers, laptops, servers
  - ARM
    - widely used on embedded, phones, recently on laptops, servers
  - RISC-V
    - fast-growing, used in embedded systems, servers, personal computers
- Even more microarchitectures
  - Apple/Samsung/Qualcomm have their own microarchitecture (implementation) of ARM
  - Intel/AMD have different microarchitectures for x86

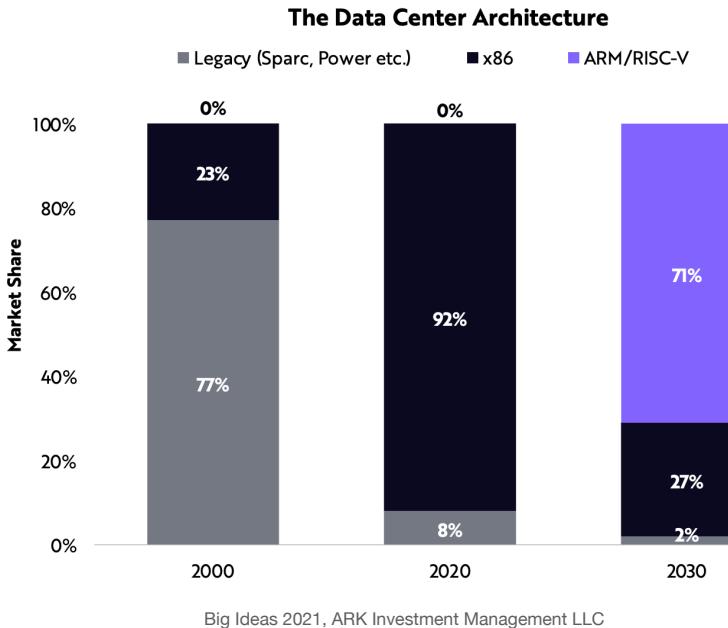
```
int sum(int *p, int n) {
    int i, sum = 0;
    for (i = 0 ; i < n ; i++) {
        sum += p[i];
    }
    return sum;
}
```



```
sum(int*, int):
    push rbp, rsp
    mov rbp, rsp
    mov qword ptr [rbp - 8], rdi
    mov dword ptr [rbp - 12], esi
    mov dword ptr [rbp - 20], 0
    mov dword ptr [rbp - 16], 0
.LBB0_1:
    mov eax, dword ptr [rbp - 16]
    cmp eax, dword ptr [rbp - 12]
    jge .LBB0_4
    mov rax, qword ptr [rbp - 8]
    movsxd rcx, dword ptr [rbp - 16]
    mov eax, dword ptr [rax + 4*rcx]
    add eax, dword ptr [rbp - 20]
    mov dword ptr [rbp - 20], eax
    mov eax, dword ptr [rbp - 16]
    add eax, dword ptr [rbp - 16]
    mov dword ptr [rbp - 16], eax
    jmp .LBB0_1
.LBB0_4:
    mov eax, dword ptr [rbp - 20]
    pop rbp
    ret
```

```
sum(int*, int):
    sub sp, sp, #16
    str r0, [sp, #12]
    str r1, [sp, #8]
    mov r0, #0
    str r0, [sp]
    str r0, [sp, #4]
    b .LBB0_1
.LBB0_1:
    ldr r0, [sp, #4]
    ldr r1, [sp, #8]
    cmp r0, r1
    bge .LBB0_4
    ldr r0, [sp]
    b .LBB0_2
.LBB0_2:
    ldr r0, [sp, #12]
    ldr r1, [sp, #4]
    ldr r0, [r0, r1, lsl #2]
    add r0, r0, r1
    str r0, [sp]
    b .LBB0_3
.LBB0_3:
    ldr r0, [sp, #4]
    add r0, r0, #1
    str r0, [sp, #4]
    b .LBB0_1
.LBB0_4:
    ldr r0, [sp]
```

```
sum(int*, int):
    addi sp, sp, -48
    sw r0, 44(sp)
    sw r1, 40(sp)
    addi a0, -48
    sw a0, -36(s0)
    sw a0, -35(s0)
    sw a1, -40(s0)
    sw zero, -24(s0)
    sw zero, -20(s0)
    j .L2
.L2:
    lw a5, -20(s0)
    lw a5, -24(s0)
    slli a5, a5, 2
    lw a4, -36(s0)
    add a5, a4, a5
    lw a4, -24(s0)
    add a5, a4, a5
    sw a5, -24(s0)
    lw a5, -20(s0)
    addi a5, a5, 1
    sw a5, -20(s0)
```



## CISC vs RISC

### Complex Instruction Set Computer (CISC)

- more complex instructions, can directly manipulate memory
- for example:
  - X86, X86\_64 (Intel and AMD, desktop/laptop/server)
  - X86\* internally are still RISC

### Reduced Instruction Set Computer (RISC)

- simpler instructions, load-store architecture
- for example:
  - ARM (smartphone/pad)
  - RISC-V** (free ISA, closer to MIPS than other ISAs)
  - Others: Power ISA, SPARC, etc

## CISC vs RISC

### CISC

- early trend was to add more and more instructions
- VAX architecture had an instruction to multiply polynomials!
- operations may directly manipulate memory

### RISC philosophy

- keep the instruction set **small and simple**
- makes it easier to build fast hardware.
- let software do complicated operations by composing simpler ones

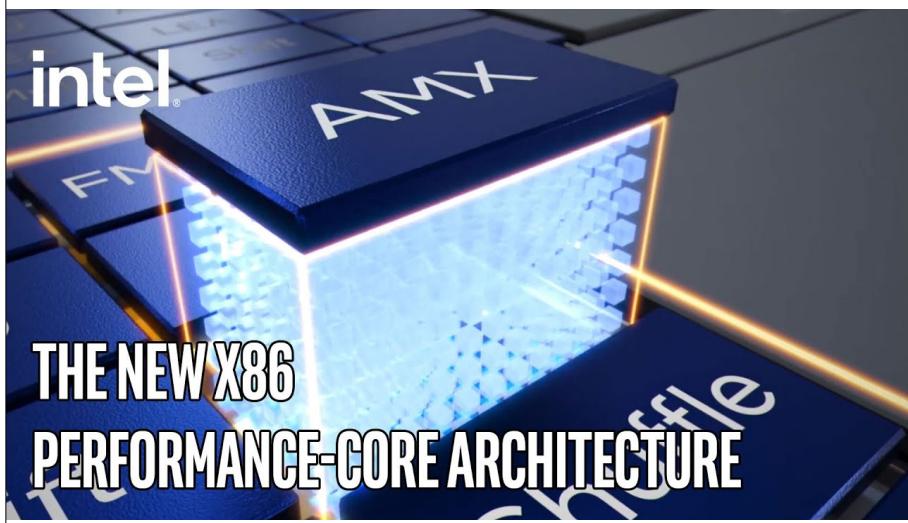


ACM has named **John L. Hennessy**, former President of Stanford University, and **David A. Patterson**, retired Professor of the University of California, Berkeley, recipients of the **2017 ACM A.M. Turing Award** for pioneering a systematic, quantitative approach to the design and evaluation of computer architectures with enduring impact on the microprocessor industry.

<https://www.youtube.com/watch?v=3LVeEjsn8Ts>

## Chip Manufacturing

### Intel's x86 core



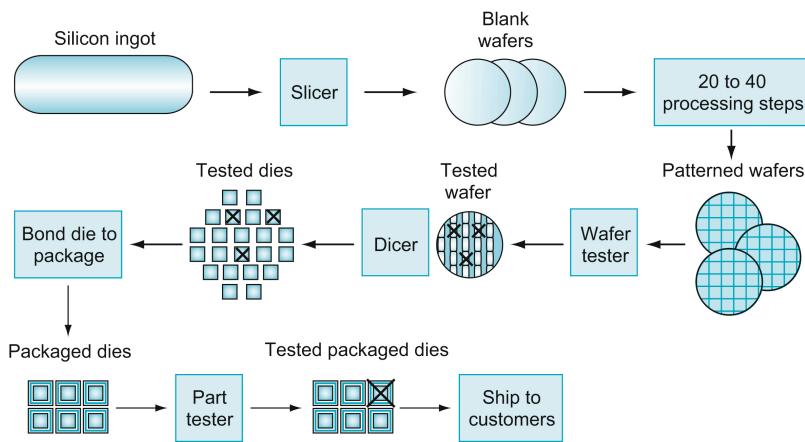
<https://www.youtube.com/watch?v=ijTRvIQV7bE>

### Semiconductor technology

- **Conductors**
  - materials that easily conduct
- **Insulators**
  - materials that do not conduct
- **Semiconductors**
  - materials with conductivity between conductors and insulators
  - those properties can be controlled by doping (adding/removing electrons)
  - Silicon (base material for computer technology)



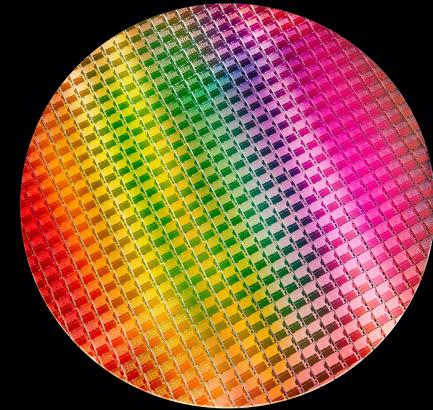
## Manufacturing integrated circuits



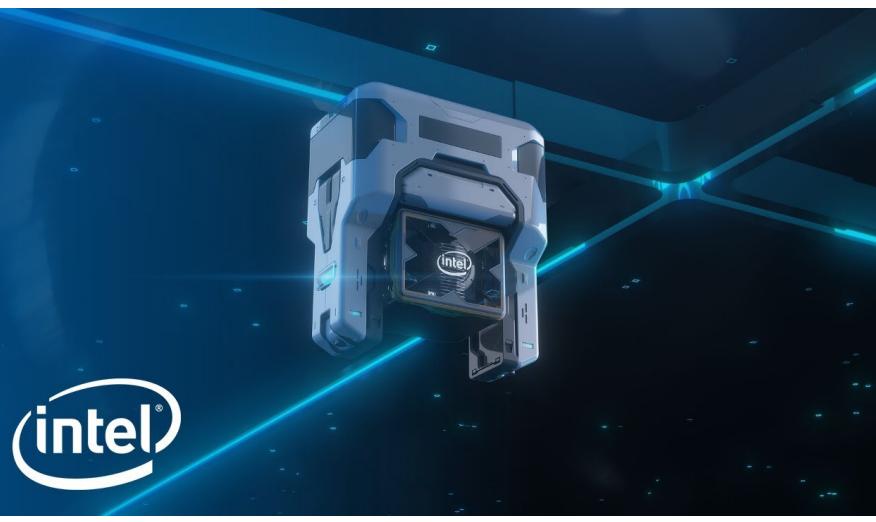
Yield: proportion of working dies per wafer

## Intel® Core 10th Gen

- 300mm wafer (almost 1 foot), 506 chips, 10nm technology
- Each chip is 11.4 x 10.7 mm (a bit less than .5 in)



## From sand to silicon



[https://www.youtube.com/watch?v= VMYPLXnd7E](https://www.youtube.com/watch?v=VMYPLXnd7E)